

Multi-core Processors



Processor development till 2004

- **Moore's Law: transistor count doubles every 18 months**
 - Used to improve processor performance by 2x every 18 months
 - Single core, binary compatible to previous generations

- **Contributors to performance improvements**
 - More ILP through OOO techniques
 - Wider issue, better prediction, better scheduling, ...
 - Better memory hierarchies
 - Clock frequency improvements

Out-of-order
Instruction scheduling



Historical Perspective

- 80's till 2004 the exponential increase in the number of transistors predicted by Moore's Law (2x /18 mons.) was used:
 - Implement sophisticated O.O.O & Superscalar designs
 - Design Large Caches
- Increasing operating frequency

Although successful for several years, this approach has eventually hit the **Power, Memory and ILP walls**

Solution: Chip Multiprocessors (CMP)


Use the transistors to increase the number of cores on a chip

New challenge: **Concurrency**



The Switch to Multi-core

The number of transistors is still increasing but more aggressive wider-issue, higher-clocked superscalars are not produced anymore, **Why??**




- ❑ **Complexity of the designs**: power, leakage, heat, design difficulty, deep pipelines
- ❑ **Power/Heat (Power Wall)**: heat, bills cooling, packaging, hot-spot, deep pipelines, cannot increase clock frequencies
- ❑ **Lack of ILP (ILP Wall)**: ILP rarely exceeds 7, with average 5
- ❑ **Marginal gain of incremental logic**:

N: transistors Perf = $O(\sqrt{N})$ Power $O(N)$



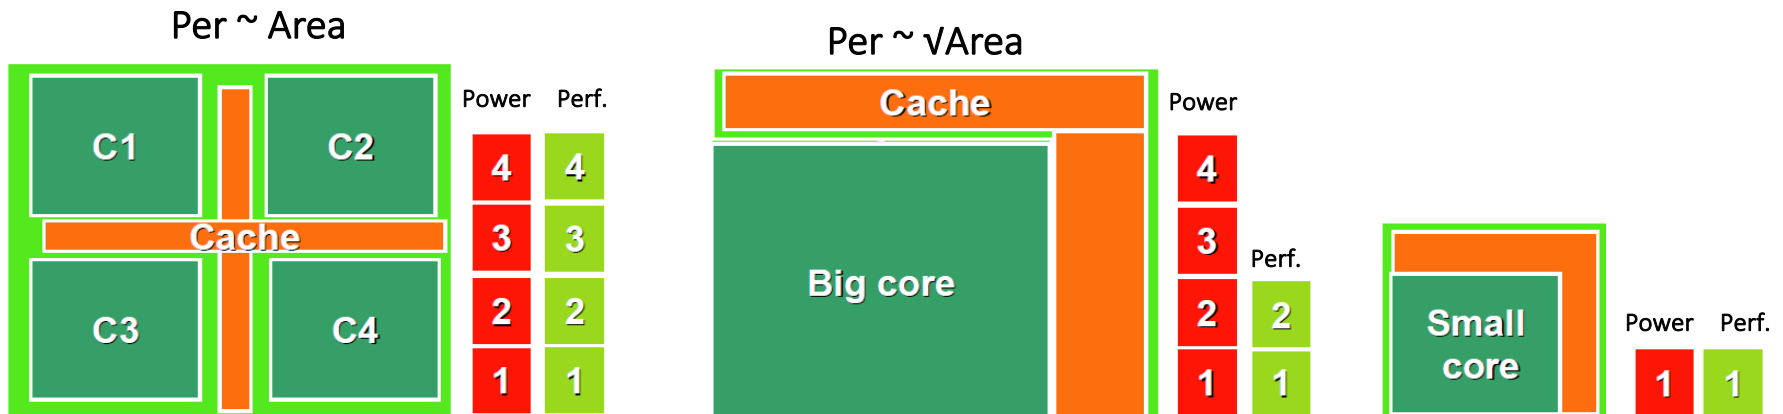
The Memory Wall





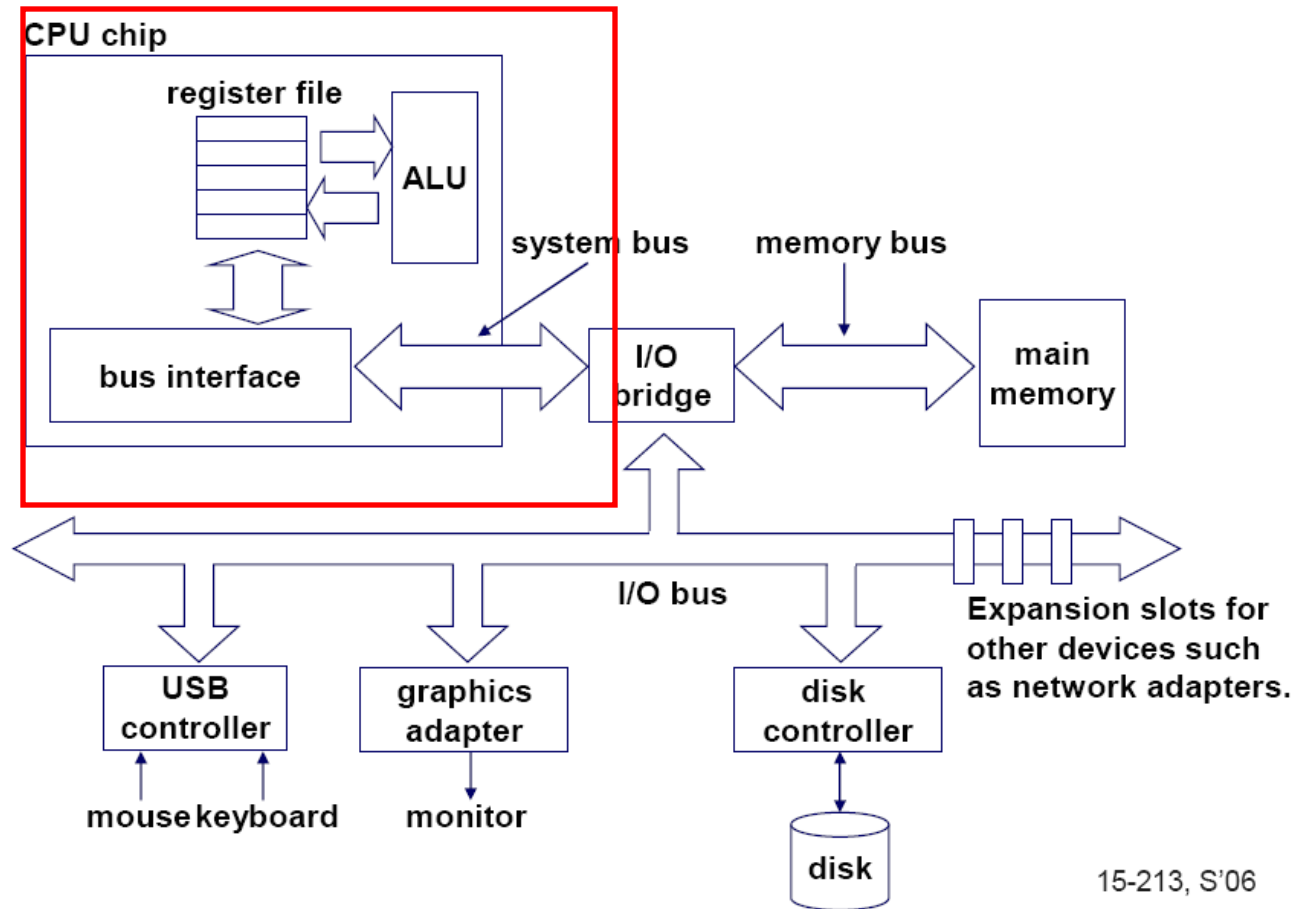
The Multi-core/Concurrency Era (Why Multi-core?)

- The solution for the previous problems is:
 - Use the silicon estate to put **more processors on the same die**
 - More area/power efficient
 - Scale the processor without complex designs
 - Shift the focus to extracting **Thread Level Parallelism (TLP)**
Currently “throughput programming”





Single-core computer

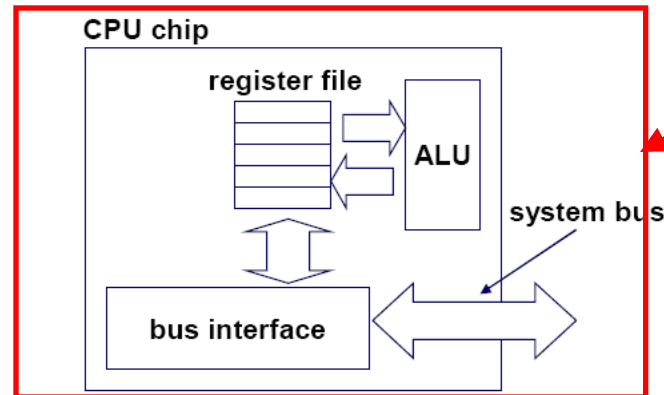


15-213, S'06



Single-core CPU chip

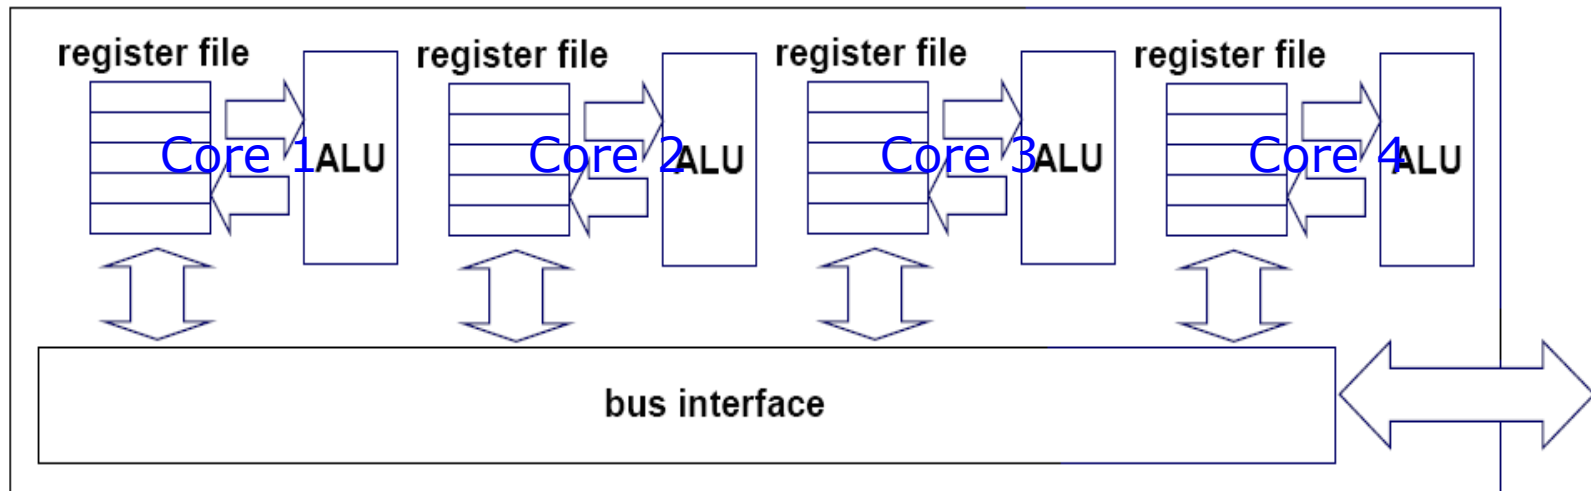
the single core





Multi-core architectures


Replicate multiple processor cores on a single die.



Multi-core CPU chip




Multi-core processor





Multi-core CPU chip

- The cores fit on a single processor socket
- Also called CMP (Chip Multi-Processor)





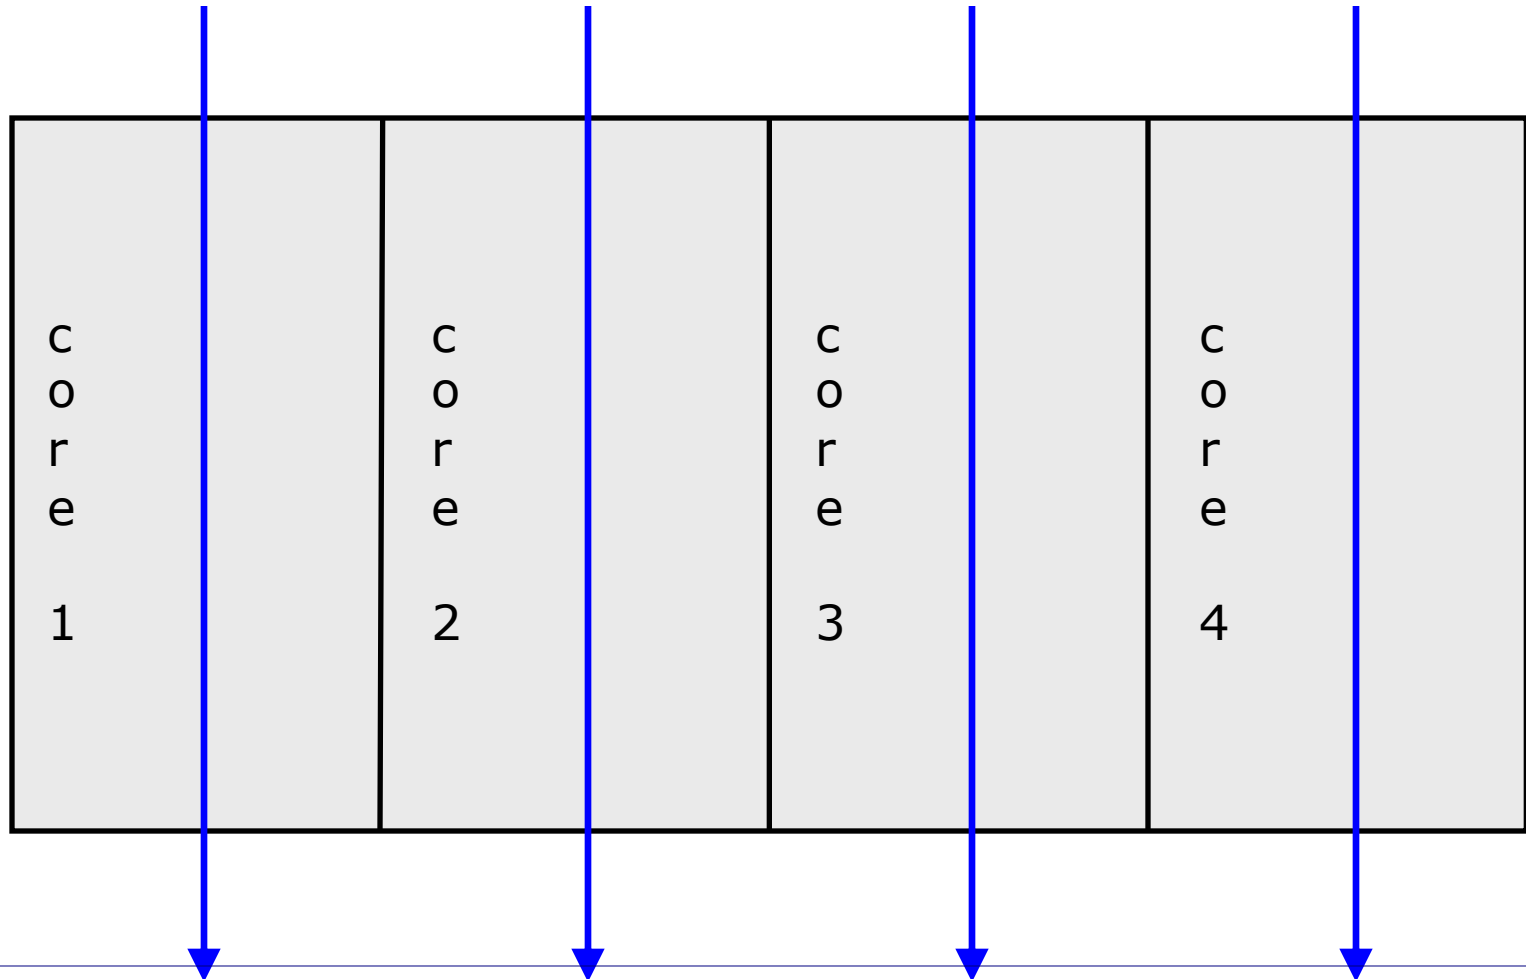
The cores run in parallel

thread 1


thread 2

thread 3

thread 4



Within each core, threads are time-sliced (just like on a uniprocessor)





Interaction with the Operating System

- OS perceives each core as a separate processor
- OS scheduler maps threads/processes to different cores
- Most major OS support multi-core today: Windows, Linux, Mac OS X, ...



Instruction-level parallelism (ILP)

- Parallelism at the machine-instruction level
- The processor can **re-order, pipeline instructions, split them** into microinstructions, do aggressive branch prediction, etc.
- **Instruction-level parallelism** enabled rapid increases in processor speeds over the last 15 years



Thread-level parallelism (TLP)

- This is parallelism on a more coarser scale
- Server can serve each client in a separate thread (Web server, database server)
- A computer game can do **AI**, **graphics**, and a **physics-related computation** in three separate threads
- Single-core superscalar processors cannot fully exploit TLP
- Multi-core architectures are the next step in processor evolution: explicitly exploiting TLP



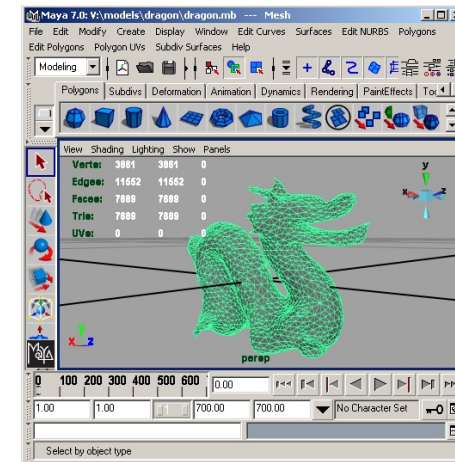
Data Level Parallelism (DLP)

- Also called SIMD or Vectorization
- More on this later



What applications benefit from multi-core?

- Database servers
- Web servers (Web commerce)
- Multimedia applications
- Scientific applications, CAD/CAM
- In general, applications with *Thread-level parallelism* (as opposed to instruction-level parallelism) are better supported



Each can run on its own core





More examples

- Editing a photo while recording a TV show through a digital video recorder
- Downloading software while running an anti-virus program
- “Anything that can be threaded today will map efficiently to multi-core”
- BUT: some applications difficult to parallelize



Multi-core design issues & Challenges

- Programming multi-cores (**Concurrency**) – is *Difficult and is the biggest challenge now*
 - ILP was **implicit** v.s. TLP and DLP **explicit**
 - Parallelizing Compilers
 - Parallel Lang./Extensions
 - H.W support
- On-chip Memory (**Memory Wall**) & Interconnection Network:
 - How to **Scale up** to tens/hundreds of cores: need on chip memory
 - Cache design/Shared vs. private Caches/**Coherency (space vs latency)**
 - **Interconnect** choices significantly affects overall performance
 - Bus/Point-to-Point/Cross-bar/Ring/Mesh/Omega (cost vs latency)
- What is the best building blocks?
 - Simple/complex/SMT/in-order/out-of-order cores?
 - Homogeneous or heterogeneous?
- Best Balance between Cores/Caches/Interconnect?




Heterogeneous Multicores




Heterogeneous vs. Homogeneous

Homogeneous Multi-core



Heterogeneous Multi-core



Big Core (Complex, out-of-order, superscalar, branch prediction, etc.)

Small Core (Simple, focus on ALU.)



Heterogeneous Multicores


- Offers advantages compared to homogeneous multicores in:
 - **Power/Area:** code requiring complex out-of-order execution with branch-predication/speculation runs on the complex cores (fast power inefficient cores).

Less control-flow code requiring ALU and data parallel (ex. SIMD/Vector) (multiple threads running on many cores)
 - **Mitigating Amdahl's Law:** execute **serial portions** on the fast and complex core and parallel portions on the smaller cores.
 - **Different types of applications run simultaneously on these cores.**




Amdahl's Law



The performance of a parallel program is limited by the serial (non-parallelizable) part of the program



Single Core








Total Time: 16s






Multi-Core (2 cores)

Total Time: 12s




Multi-Core (4 cores)

Total Time: 10s



...



Multi-Core (∞ cores)

Total Time: 8s



Thread Level Parallelism (TLP)



SIMD

Vectorization

DLP




Motivation

- Intel analyzed multimedia applications and found they share the following characteristics:
 - Small native data types (8-bit pixel, 16-bit audio)
 - Recurring operations (same instruction on many pieces of large data sets).
 - Inherent parallelism
- This gave birth to Vector/SIMD support at the level of instructions in mainstream-processors.

Scalar vs. SIMD

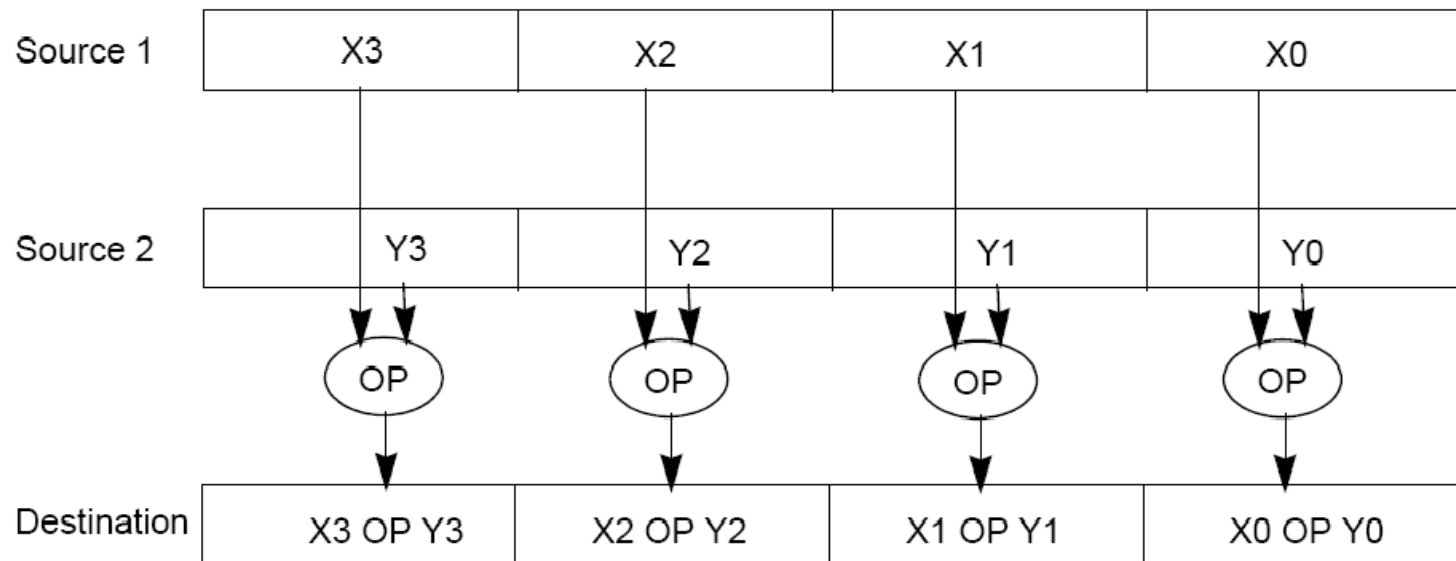
- **Scalar processing**
 - **traditional mode**
 - **one operation produces one result**





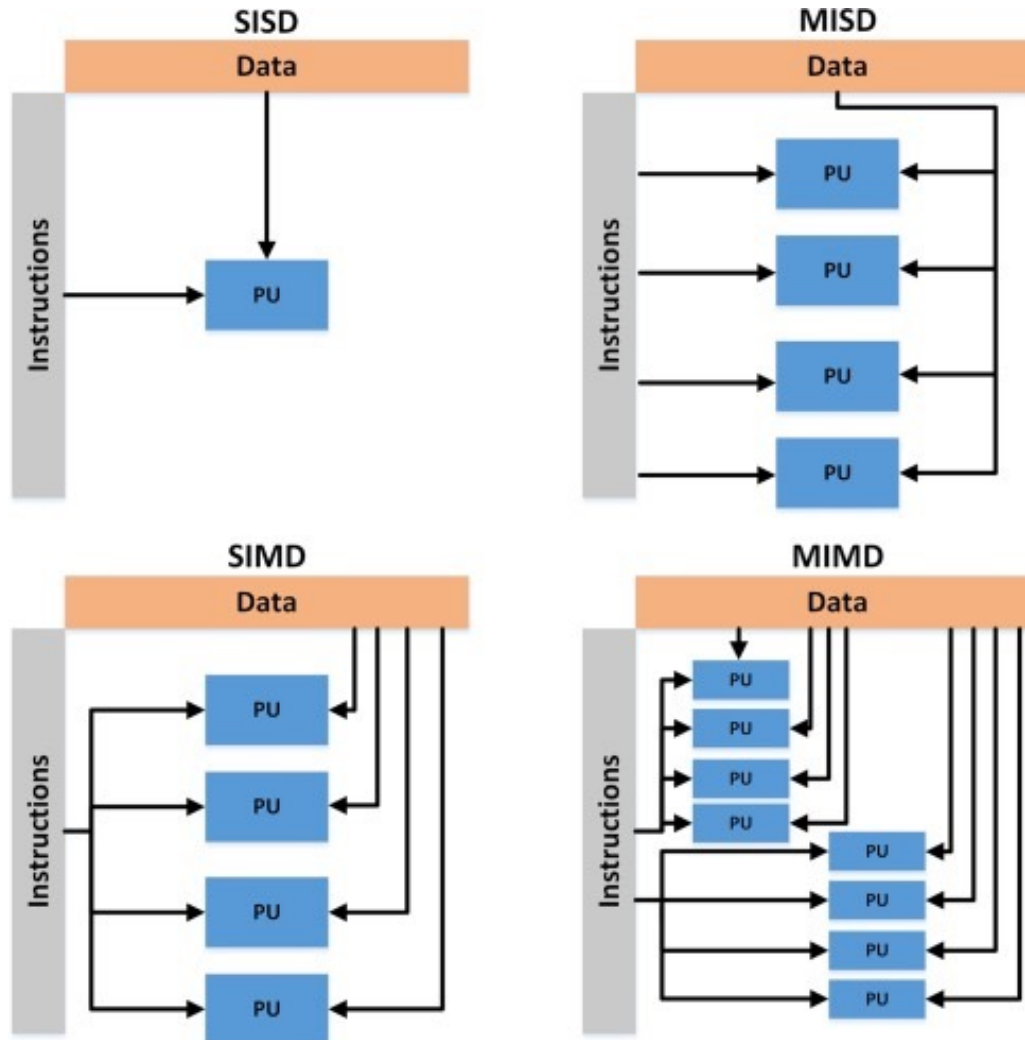
SIMD

- SIMD (single instruction multiple data) architecture performs the same operation on multiple data elements in parallel
- **PADDW MM0, MM1**





Flynn Taxonomy





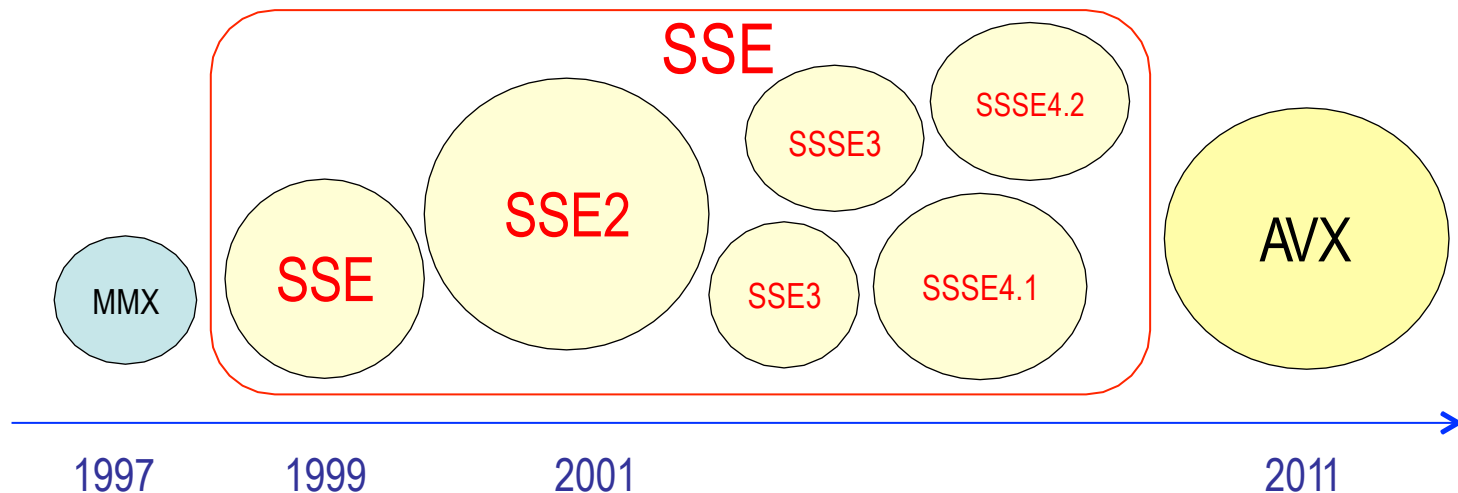
SIMD Architecture Support

- **MMX** (Multimedia Extension) was introduced in 1996 - P1 / PII.
 - 64-bit vectors / Only integer operations
- **SSE** (Steaming SIMD Extension) - Pentium III.
 - 128-bit vectors (only floating-point)
- **SSE2** - P4. (floating-point, double & Integer)
- **SSE3** - P4/HT
 - 13 more instructions.
- **AVX** (Advanced Vector Extensions) – 2011
 - 256-bit vectors
- SIMD extensions for other architectures:
 - 3DNow! / AltiVec / VIS / VMX



The SSE family of Vector extensions


- Include over 400 instructions
- Available in most modern processors
- Uses 16 dedicated registers of length 128 bits
- Programming with AVX is similar, vector length and instruction names differ






MMX data types


Packed Byte: 8 bytes packed into 64 bits



Packed Word: 4 words packed into 64 bits



Packed Doubleword: 2 doublewords packed into 64 bits




Packed Quadword: One 64-bit quantity







Application: frame difference






Application: frame difference

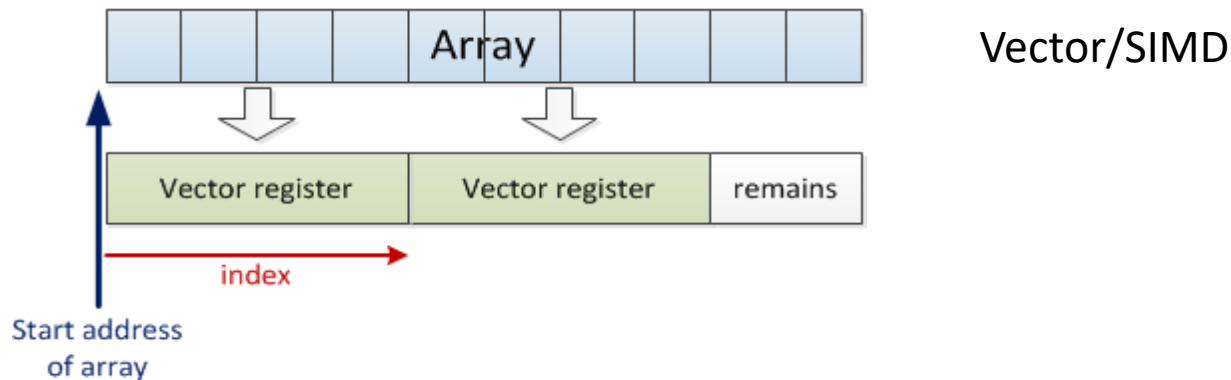
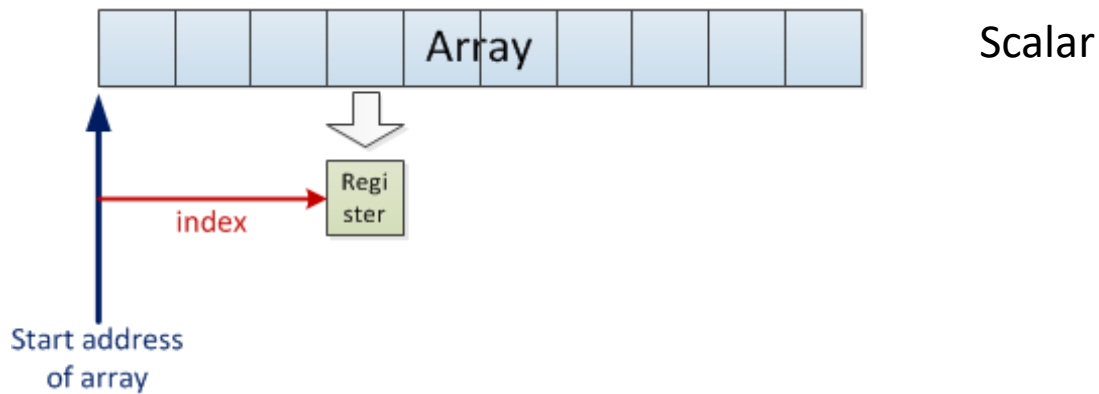


(A-B) or (B-A)





Loading from array to register





Application: frame difference

```
MOVQ      mm1, A //move 8 pixels of image A
MOVQ      mm2, B //move 8 pixels of image B
MOVQ      mm3, mm1 // mm3=A
PSUBSB    mm1, mm2 // mm1=A-B
PSUBSB    mm2, mm3 // mm2=B-A
POR       mm1, mm2 // mm1=|A-B|
```




Performance boost (data from 1996)

Benchmark kernels:
FFT, FIR, vector dot-product, IDCT,
motion compensation.

65% performance gain

Lower the cost of
multimedia programs

by removing the need
of specialized DSP
chips





Issues with MMX

- ❑ Only supported integer operations
- ❑ MMX and FPU can not be used at the same time.
 - Used the same set of registers
 - Big overhead to switch.
- ❑ This proved to be a bad decision later.
- ❑ It is why Intel introduced SSE later as a separate unit.

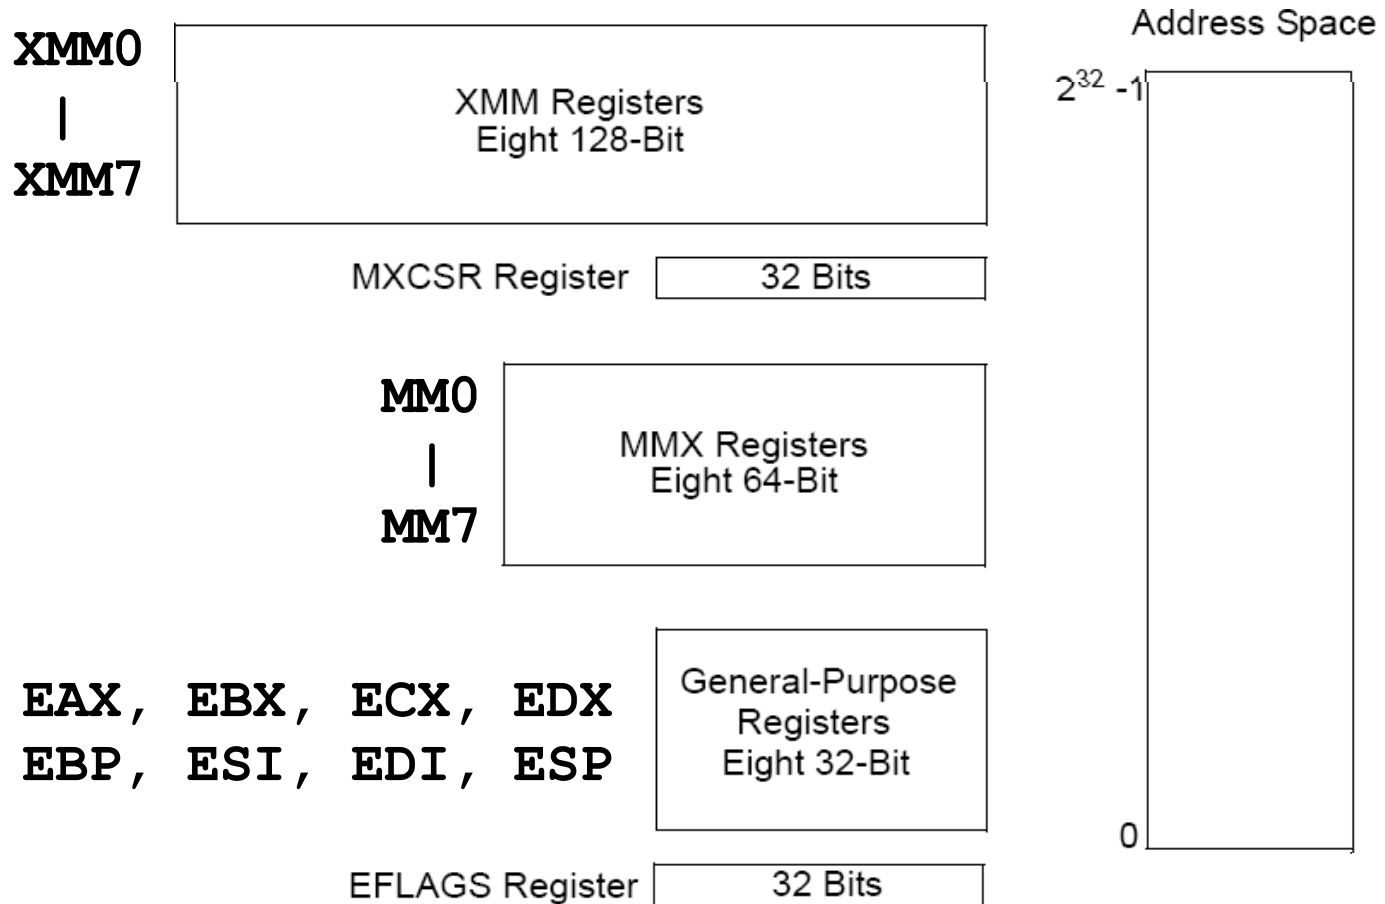


SSE

- Adds eight 128-bit registers
- Allows SIMD operations on packed single-precision floating-point numbers
 - Later added support for double-precision FP
- Most SSE instructions require 16-aligned addresses
- Instructions that explicitly prefetch data, control data cacheability and ordering of store

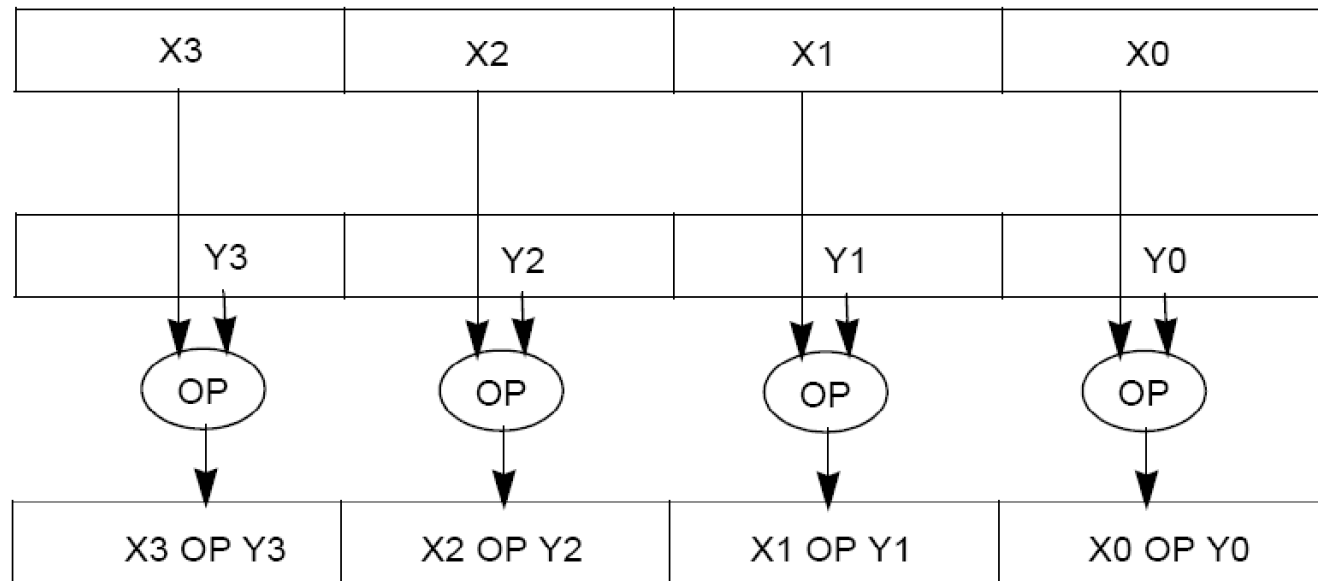


SSE programming environment






SSE packed FP operation



- **ADDPS/SUBPS:** packed single-precision FP



SSE scalar FP operation








- **ADDSS/SUBSS**: scalar single-precision FP



SSE features

- Add data types and instructions for them



Programming environment unchanged



How to use assembly in projects

- ☐ Write the whole project in assembly
 - Difficult/rare
- ☐ Link with high-level languages
 - Develop most of the program in HLL
 - Use assembly for performance critical parts
- ☐ Inline assembly
 - Inline Assembly code directly into a HLL program.
 - Compilers such as Visual C++ have compiler-specific directives to identify inline ASM code.
- ☐ Intrinsics



Intrinsics

- ❑ A function known by the compiler that directly maps to a sequence of assembly instructions.
- ❑ The compiler manages :
 - Register names
 - Register allocations
 - Memory locations of data
- ❑ More efficient than called functions
 - No calling linkage is required.

`_mm_<opcode>_<suffix>`

ps: packed single-precision

ss: scalar single-precision



Intrinsics

```
#include <xmmintrin.h>
```

```
__m128 a , b , c;  
c = _mm_add_ps( a , b );
```

```
float a[4] , b[4] , c[4];  
for( int i = 0 ; i < 4 ; ++ i  
    )  c[i] = a[i] + b[i];
```

```
// a = b * c + d / e;
```

```
__m128 a = _mm_add_ps( _mm_mul_ps( b , c ) ,  
                      _mm_div_ps( d , e ) );
```



More Examples



VMX Example (PowerPC)

E.g. 1: Use in variable initialization statement (Vector literal is bolded)

```
__vector signed int va = (__vector signed int) { -2, -1, 1, 2 };  
va = vec_add(va, ((__vector signed int) { 1, 2, 3, 4 }));
```

E.g. 2: Accessing vector as a scalar - the third element of a vector (Pointer cast is bolded)

```
__vector signed int va = (__vector signed int) { 1, 2, 3, 4 };  
int *a = (int *) &va;  
printf("a[2] = %d", a[2]);
```

E.g. 3: Accessing a scalar array as a vectors (Pointer cast is bolded)

```
int a[8] __attribute__((aligned(16))) = { 1, 2, 3, 4, 5, 6, 7, 8 };  
__vector signed int *va = (__vector signed int *) a;  
  
// va[0] = { 1, 2, 3, 4 }, va[1] = { 5, 6, 7, 8 }  
vb = vec_add(va[0], va[1]);
```



VMX Example (PowerPC)

```
#include <stdio.h>
#include <altivec.h>

// declares input/output scalar variables

int a[4] __attribute__((aligned(16))) = { 1, 3, 5, 7 };
int b[4] __attribute__((aligned(16))) = { 2, 4, 6, 8 };
int c[4] __attribute__((aligned(16)));

int main(int argc, char **argv)
{
    // declares vector variables which points to scalar arrays

    __vector signed int *va = (__vector signed int *) a;
    __vector signed int *vb = (__vector signed int *) b;
    __vector signed int *vc = (__vector signed int *) c;

    // adds four signed integers at once

    *vc = vec_add(*va, *vb);    // 1 + 2, 3 + 4, 5 + 6, 7 + 8

    // output results

    printf("c[0]=%d, c[1]=%d, c[2]=%d, c[3]=%d\n", c[0], c[1], c[2], c[3]);

    return 0;
}
```



SSE Example (intel)

```
float Scalar(float *s1, float *s2) {  
    int i;  
    float prod;  
  
    for(i=0; i<size; i++) {  
        prod += s1[i] * s2[i];  
    }  
    return prod;  
}
```

```
float ScalarSSE(float *s1, float *s2) {
```

```
    float prod;
```

```
    int i;
```

```
    __m128 X, Y, Z;
```

```
    for(i=0; i<size; i+=4) {  
        X = _mm_load_ps(&s1[i]);  
        Y = _mm_load_ps(&s2[i]);  
        X = _mm_mul_ps(X, Y);  
        Z = _mm_add_ps(X, Z);  
    }
```

```
    for(i=0; i<4; i++) {  
        prod += Z[i];  
    }  
    return prod;
```

```
}
```

$Z[0] = s1[0] * s2[0] + s1[4] * s2[4] + s1[8] * s2[8] + \dots$

$Z[1] = s1[1] * s2[1] + s1[5] * s2[5] + s1[9] * s2[9] + \dots$


$Z[2] = s1[2] * s2[2] + s1[6] * s2[6] + s1[10] * s2[10] + \dots$

$Z[3] = s1[3] * s2[3] + s1[7] * s2[7] + s1[11] * s2[11] + \dots$



Other SIMD architectures

- Graphics Processing Unit (GPU)





References

- *Intel MMX for Multimedia PCs*, CACM, Jan. 1997
- Chapter 11 *The MMX Instruction Set*, The Art of Assembly
- Chap. 9, 10, 11 of IA-32 Intel Architecture Software Developer's Manual: Volume 1: Basic Architecture
- http://www.csie.ntu.edu.tw/~r89004/hive/sse/page_1.html